

# Designing Indexes to Improve Query Performance: Part 2

---



**Gail Shaw**

TECHNICAL LEAD

@SQLintheWild <http://sqlinthewild.co.za>



# Overview



Introducing nonclustered indexes

Common query predicates

Indexing for equality

Indexing for inequality

Indexing for ORs

Indexing for joins

Include columns

Filtered indexes



# Indexing for Joins

Temp Icons! Replace in editing



Nested loop joins benefit from an index on the inner table



Merge joins may benefit from indexes to provide necessary ordering



Hash joins don't benefit from indexes



# Demo



## Indexing for joins



# Index Include Columns

**Additional columns at the leaf level  
of the index**

**Used to avoid expensive key lookups**

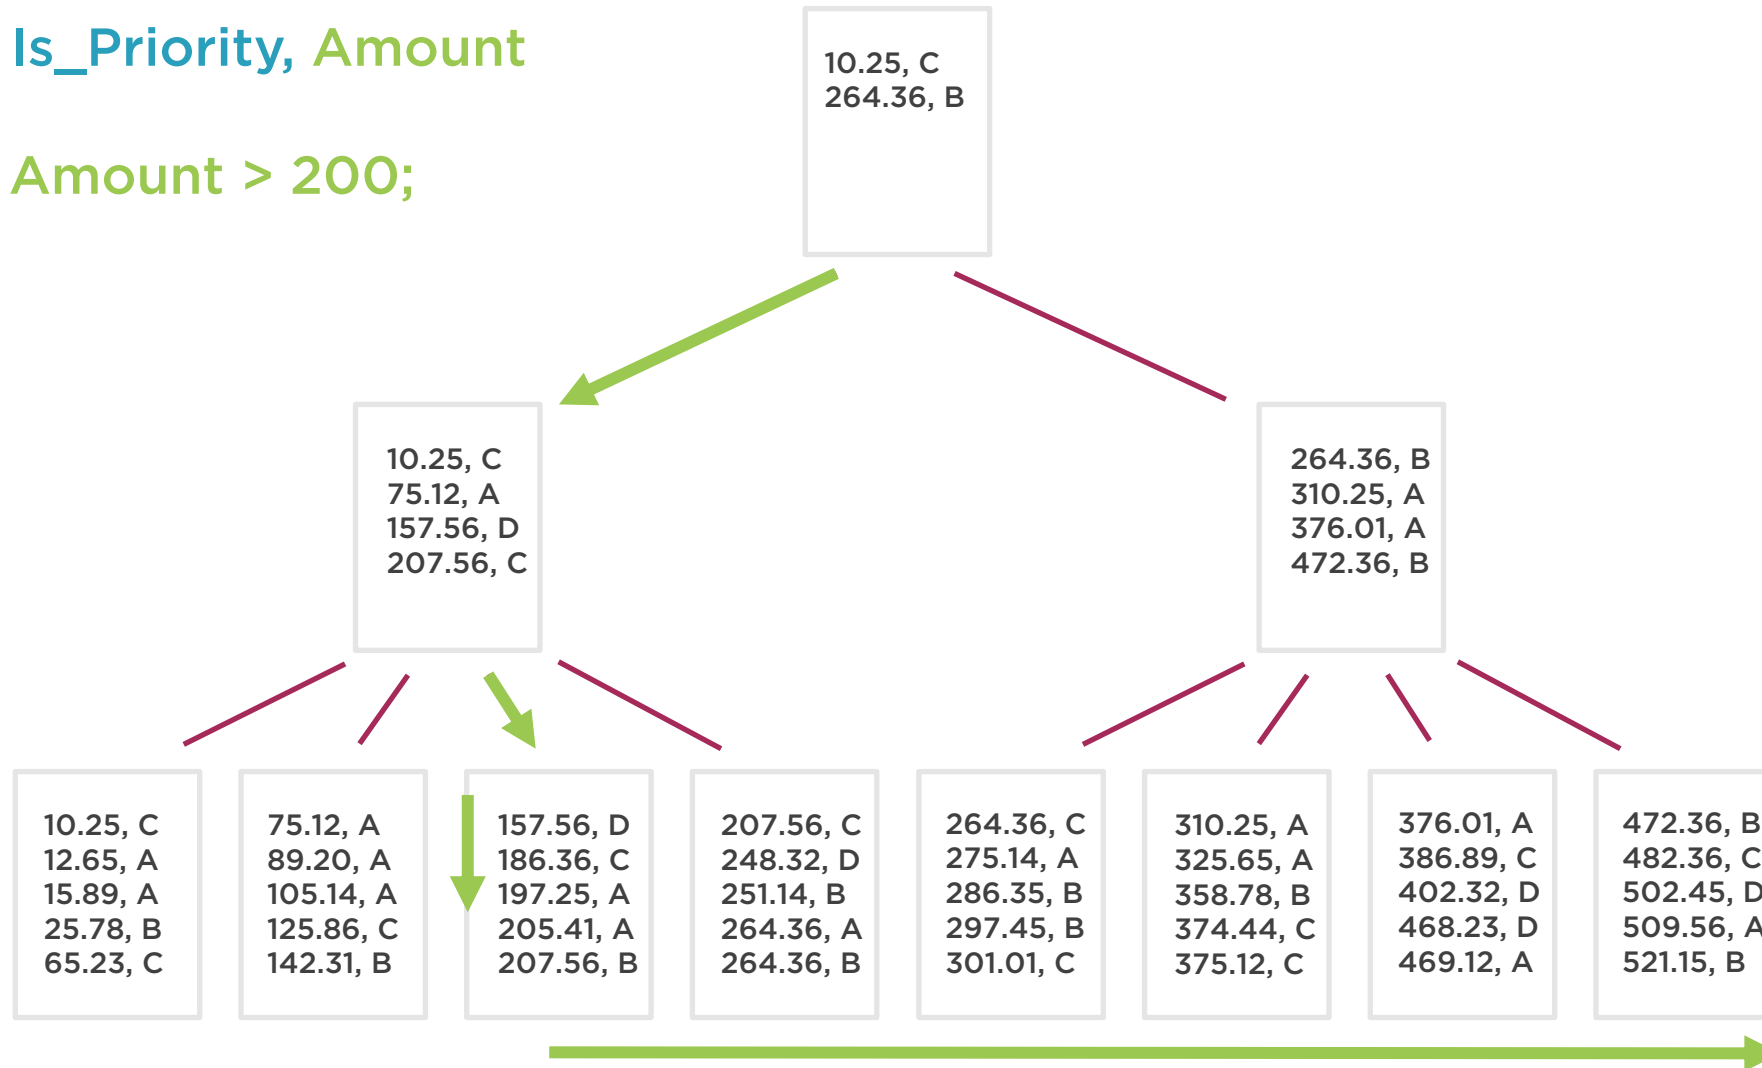


# Include Columns in an Index

**SELECT** **Is\_Priority**, **Amount**

...

**WHERE** **Amount** > 200;



# Key Lookups

Single-row seek  
against the  
clustered index

Fetch columns  
which are  
required, but not  
in the key of the  
index used

Slow if there are a  
large number




# Include Columns in an Index

SELECT Is\_Priority

...

WHERE Amount > 200;





# Demo



Include columns



# Filtered Indexes

**Indexes on a subset of rows in the table**

**Can be useful on tables with skewed data**

**Also useful for complex unique constraints**

**Don't work with parameterised queries**



# Demo



## Filtered indexes



# How Many Nonclustered Indexes?



**As many as you need for the  
workload**



**And no more**

# Summary



Introducing nonclustered indexes

Indexing for equality

Indexing for inequality

Indexing for ORs

Indexing for joins

Include columns

Filtered indexes

