

Designing Columnstore Indexes for Analytic Queries



Gail Shaw

TECHNICAL LEAD

@SQLintheWild <http://sqlinthewild.co.za>



Agenda



Introducing Columnstore indexes

Columnstore architecture

Clustered and nonclustered Columnstore indexes

When to use Columnstore indexes



Columnstore Indexes

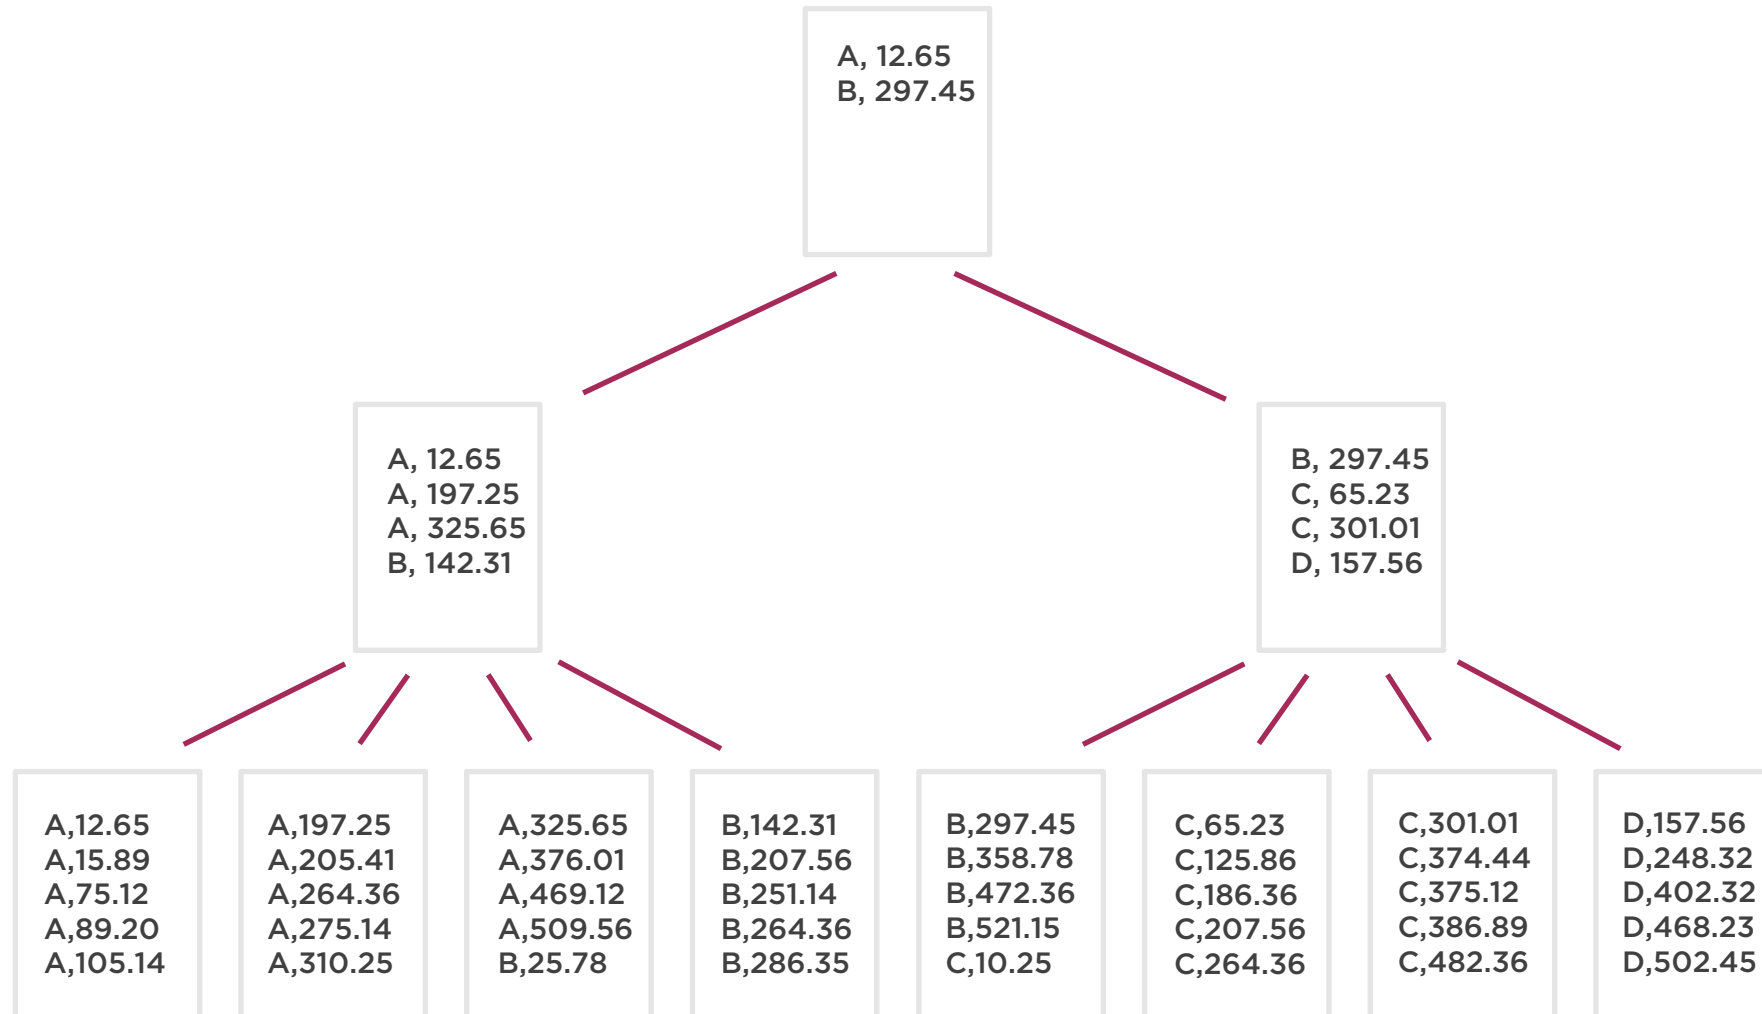
Indexes where the data for a column is stored together

Highly compressed

Good for analytics-type queries



Rowstore Index Architecture



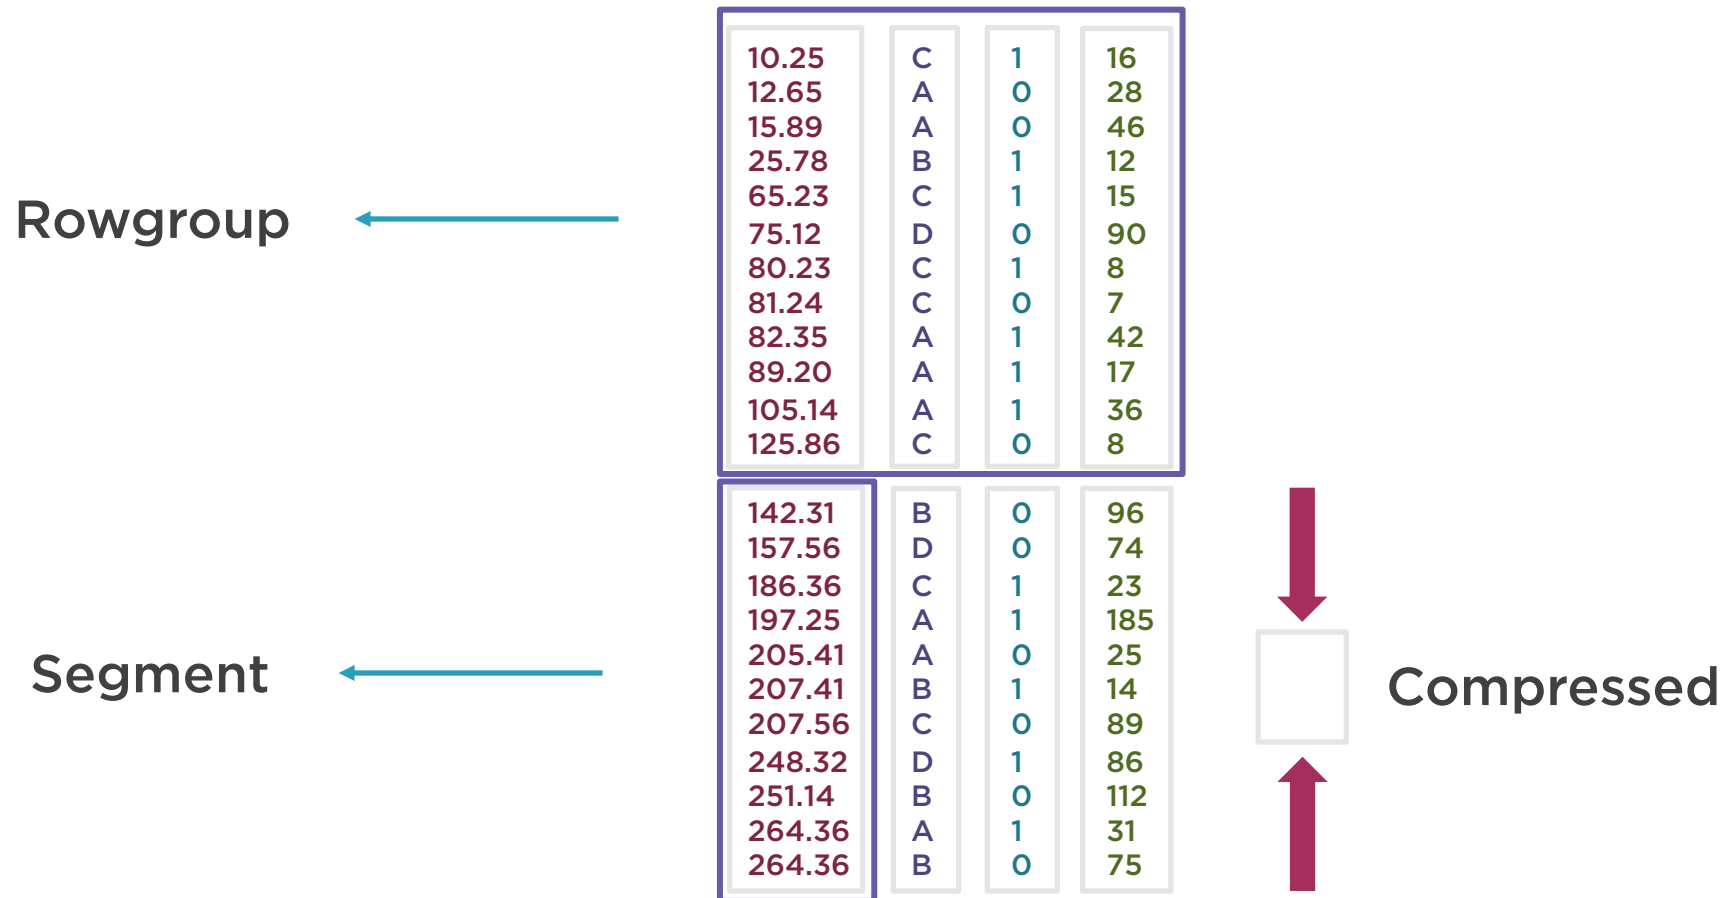
Columnstore Index Architecture

CREATE COLUMNSTORE INDEX
ON Transactions
(Amount, TransactionType,
IsPriority, ShipmentID)


10.25	C	1	16
12.65	A	0	28
15.89	A	0	46
25.78	B	1	12
65.23	C	1	15
75.12	D	0	90
80.23	C	1	8
81.24	C	0	7
82.35	A	1	42
89.20	A	1	17
105.14	A	1	36
125.86	C	0	8
142.31	B	0	96
157.56	D	0	74
186.36	C	1	23
197.25	A	1	185
205.41	A	0	25
207.41	B	1	14
207.56	C	0	89
248.32	D	1	86
251.14	B	0	112
264.36	A	1	31
264.36	B	0	75




Columnstore Index Architecture



Modifying Data in a ColumnStore Index



Deleted List



Deltastore

10.25	C	1	16
12.65	A	0	28
15.89	A	0	46
25.78	B	1	12
65.23	C	1	15
75.12	D	0	90
80.23	C	1	8
81.24	C	0	7
82.35	A	1	42
89.20	A	1	17
105.14	A	1	36
125.86	C	0	8
142.31	B	0	96
157.56	D	0	74
186.36	C	1	23
197.25	A	1	185
205.41	A	0	25
207.41	B	1	14
207.56	C	0	89
248.32	D	1	86
251.14	B	0	112
264.36	A	1	31
264.36	B	0	75

Types of Columnstore Indexes

Clustered columnstore

Organises the table, consists of all the columns in the table

Nonclustered columnstore

Secondary structure, contains just the columns specified



Demo



Creating and examining columnstore Indexes



Columnstore Guidelines

**Clustered
columnstore
indexes on large
fact tables**

**Can add
nonclustered
rowstore indexes
to help with
searches**

**Nonclustered
columnstore
indexes can help
with analytics in
OLTP systems**



Considerations for Columnstore Indexes

Not ideal for
systems with lots of
single-row
modifications

Not efficient on small
tables. Rowgroups
should ideally have ~ 1
million rows

Compression works
best on columns
with repeated
patterns in the data

Can't be used for
seek operations

Can't be used to
enforce
uniqueness



Order of Columns

Not of major importance with a Columnstore

No 'left-based subset of the key'

For nonclustered Columnstores, ensure they contain the columns to be aggregated and grouped



Demo



Query Performance with Columnstore Indexes



Summary



Columnstore index architecture

Guidelines for creating columnstore indexes

